

# Convolutional Neural Networks

Julius Taylor

0000000

stud3@email.com

Shawn Cala

4921431

shawn.cala@gmail.com


February 5, 2017

# Table of Contents


- 1 Background
  - Neural Networks
  - Structure
- 2 Convolutional Neural Networks
  - Concept
  - Kernel function
  - Application
- 3 Components
  - Activation function
  - Backpropagation
- 4 Problems
  - Overfitting
  - Underfitting
- 5 Code

# Neural Networks

Primary motivation: Neural Networks mathematically simulate biological functionalities of the human brain



(a) biological model



(b) mathematical representation

Figure: neuronal model and computational abstraction

# Structure

Neural Networks generally contain:

- an n-dimensional input
- one or many layers of interconnected neurons
- an output-layer




Figure: Basic concept of a Neural Network

# The Problem Space

- image classification: Image  $\rightarrow$  Class that describes the image (eg. 'Cat')
- recognizing patterns and generalizing from prior knowledge
- figure out features, that describe things

# Problems with image classification

- naive approach: comparing pixels one by one
- very sensitive to small changes and does not generalize



# Problems with image classification

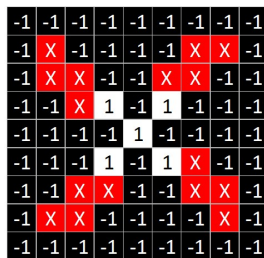




Figure: Visualizing what computers see

# Problems with image classification

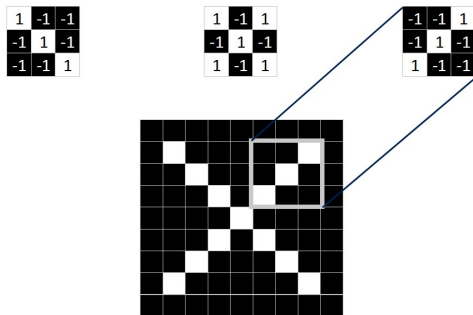




# What makes us conclude that two pictures match?



# What makes us conclude that two pictures match?

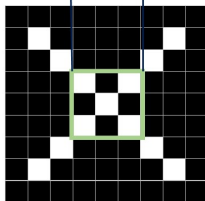


# What makes us conclude that two pictures match?

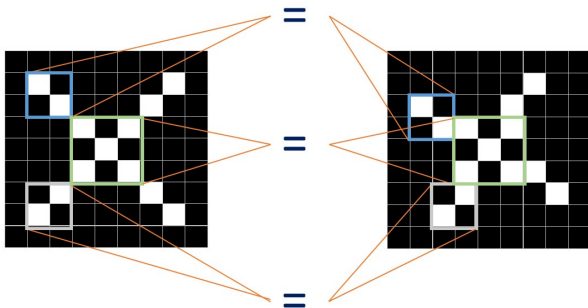
1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1

-1	-1	1
-1	1	-1
1	-1	-1



# What makes us conclude that two pictures match?



Convolutional Neural Networks (CNNs) are a subtype of Neural Networks:

- is a type of feed-forward network
- structure inspired by the animal visual cortex
- uses many identical copies of the same neuron
- express computationally extensive models while keeping weights to learn small
- breakthrough results in pattern recognition problems

# Structure

Convolution



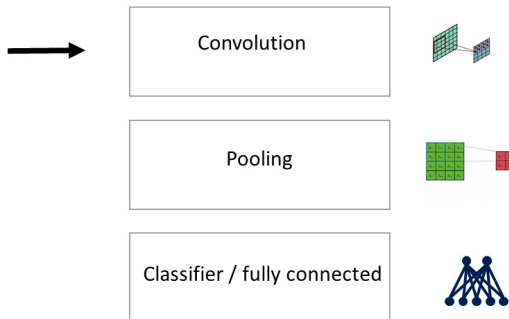
Pooling



Classifier / fully connected




# Structure



# Finding Features

0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Pixel representation of filter




Visualization of a curve detector filter



# Finding Features



Original image



Visualization of the filter on the image

# Finding Features



Visualization of the  
receptive field

0	0	0	0	0	0	30
0	0	0	0	50	50	50
0	0	0	20	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0

Pixel representation of the receptive  
field


\*

0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Pixel representation of filter


Multiplication and Summation =  $(50*30)+(50*30)+(50*30)+(20*30)+(50*30) = 6600$  (A large number!)

# Finding Features



Multiplication and Summation = 0

# Finding Features



$$(f * g)(c_1, c_2) = \sum_{a_1, a_2} f(a_1, a_2) \cdot g(c_1 - a_1, c_2 - a_2)$$

# First Trick: Convolutions


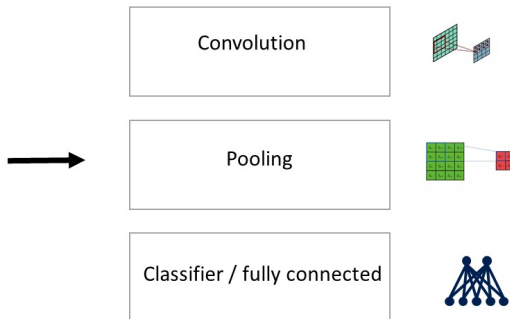


Figure: A convolutional layer fed into a fully connected layer

# Structure



# Second Trick: Pooling





Figure: Added max pooling to previous net

# Second Trick: Pooling


- 'zooms' out
- small patch after pooling corresponds to much larger patch before it
- makes the network a little more invariant to location of features





# Stacking Layers

- layers are composable
- we can feed the output of one layer into the input of another
- with each layer, we can detect higher level, more abstract features



# Stacking Layers


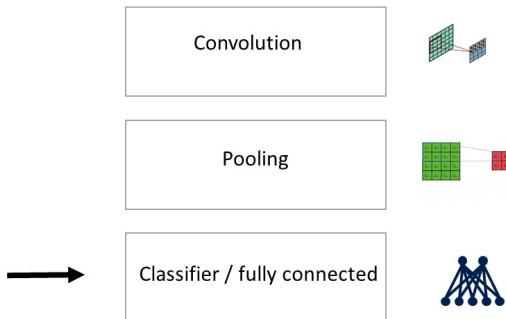


Figure: Visualizing a ConvNet trained on handwritten digits

# Structure



# Third Trick: Fully Connected

- every value in a fully connected layer gets a 'vote' which affects the result of classification
- values in certain positions are associated with a certain class




Figure: Weights in a fully connected layer contributing to classification

# Application

## Visual Object Recognition



**Figure:** ImageNet Classification with Deep Convolutional Neural Networks

ImageNet by Krizhevsky et al (2012) classified 1.2 million high-resolution images in the ImageNet LSVRC-2010 into 1000 different classes. It achieves error rates of 37.5% for the top result and 17.0% for the top-5 results

# Application

## Text Classification




Figure: Excerpt of mnist

fehlt: etwas Erklärung hierzu

# Components

hier Bild eines CNNs unseres Typs mit Komponenten

- 1 Neuron
- 2 Kernel-Function
- 3 Sigmoid function
- 4 softmax

# Activation Function

Typically, an activation function is used to add nonlinearity to the network




Figure: Different activation functions

Activation functions must be non linear and differentiable.  
The employed network uses Sigmoid:

$$\sigma(x) = \frac{1}{(1+e^{(-x)})}$$



# Components - Neuron in fully connected layer




Figure: Single neuron of a fully connected layer

Input: connected neurons from the previous layer


Output:  $f(\sum_i w_i x_i)$

with

$w_i$ :

# Components - Neuron in convolutional layer

To understand the concept behind Convolution we first demonstrate it in an example:



**Figure:** Probability distribution of ball thrown twice

We want to calculate the likelihood of the ball traveling a distance  $c$  after being thrown twice in a row.

## Input layer

The input layer consists of an  $n$ -dimensional matrix, which contain the information to be processed

- Our example uses a 2-dimensional input which represents the pixels of the images
- asd

Input:

$a, b$ : length of first and second throw  $f(a), g(b)$ : probability distribution of the balls location If we want  $c = a + b$  the probability of this event is  $f(a) \cdot g(b)$


Since only the actual outcome is relevant to us, the actual values of  $a$  and  $b$  can be varied.

The total likelihood of the ball landing at  $c$  can also be summed over all probability distributions where  $a + b = c$ :

$$f * g(c) = \sum_{a+b=c} f(a) \cdot g(a)$$

The result of this sum is a convolution between  $a$  and  $b$

Image processing with a kernel function can also be described as a



convolution:

**Figure:** Convolution between an input image and a kernel function

$$out = \sum_a \sum_b w_{ab} \cdot x$$

# Backpropagation

Backpropagation is the key algorithm which makes training Neural Networks feasible by greatly increasing the efficiency of learning algorithms.

To understand Backpropagation we first introduce the concept of forward propagation:

# Forward propagation

Forward propagation is the process of computing the output of a network for a given input




Figure: Computational graph of  $e = (a + b) \cdot (b + 1)$

# Forward propagation

We now arbitrarily set  $a = 2$  and  $b = 1$  for our example:




Figure: Computational graph of  $e = (a + b) \cdot (b + 1)$



# Forward propagation

If we want to train our network weights we must know how the output of the CNN is influenced by its layer weights. In our example our interest is how the result  $e$  is influenced by  $a$  and  $b$ . This is achieved by finding the respective partial derivatives





Figure: Partial derivatives

$c$  is only indirectly influenced by  $a$  and  $b$ . To derive the exact influence, the chain rule can be applied:  $\frac{\partial e}{\partial a} = \frac{\partial e}{\partial c} \cdot \frac{\partial c}{\partial a}$

$$\frac{\partial e}{\partial b} = \frac{\partial e}{\partial d} \cdot \frac{\partial e}{\partial c} \cdot \frac{\partial d}{\partial b} \cdot \frac{\partial c}{\partial b}$$


# Forward propagation

In principle, this approach works for training a network. There are, however, obvious problems with exponential amounts of calculations if too many possible paths between input and output exist.




# Backpropagation

Instead of determining the influence of one input on the output we calculate how the output is influenced by the previous layer.



# Backpropagation


In our previous example executing a backward differentiation immediately gives us the derivative of the output relative to every input.



= massive speed in networks with many inputs/layers

# Backpropagation

In our previous example executing a backward differentiation immediately gives us the derivative of the output relative to every input.



= massive speed in networks with many inputs/layers

# Backpropagation

Lets reiterate our network:

Convolution:  $conv = x \cdot W_1 + b_1$

Convolutional layer output:  $out = \tanh(conv)$

Fully connected layer output:  $\hat{y} = out \cdot W_2 + b_2$

Loss function:  $L = - \sum_N y \cdot \ln \hat{y}$

$$\frac{\partial L}{\partial W_2} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial W_2}$$

$$\frac{\partial L}{\partial \hat{y}} = \sum_N y \cdot \frac{\partial \ln \hat{y}}{\partial \hat{y}} = \hat{y} - y$$

$$\frac{\partial \hat{y}}{\partial W_2} = \frac{\partial out \cdot W_2 + b_2 = out^T}{\partial W_2} \frac{\partial L}{\partial W_2} = (\hat{y} - y) * out^T$$

$$\frac{\partial L}{\partial b_2} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial b_2}$$

$$= (\hat{y} - y) \cdot 1$$

$$\frac{\partial L}{\partial W_1} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial out} \cdot \frac{\partial out}{\partial conv} \cdot \frac{\partial conv}{\partial W_1}$$

$$\frac{\partial \hat{y}}{\partial conv} = \frac{\partial out \cdot W_2 + b_2}{\partial out} = W_2 \frac{\partial out}{\partial conv} = \frac{\partial \tanh(conv)}{\partial conv} = 1 - \tanh^2 conv$$

$$\frac{\partial out}{\partial W_1} = x^T$$

$$\frac{\partial L}{\partial W_1} = x^T \cdot (1 - \tanh^2 conv \cdot (y - \hat{y}) \cdot W_2$$

$$\frac{\partial L}{\partial b_1} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial out} \cdot \frac{\partial out}{\partial conv} \cdot \frac{\partial conv}{\partial b_1} =$$

$$= (1 - \tanh^2 conv \cdot (y - \hat{y}) \cdot W_2 \cdot 1$$

# Overfitting

Overfitting occurs when the complexity of the model relative to the training size is too high.

The model begins memorizing the training data rather than the underlying principle and easily loses its predictive power when the input is slightly altered.

Overfitting is prevented in a number of ways:

- 1 max-pooling layers
- 2 dropout layers
- 3





# Code



Leslie Lamport, *L<sup>A</sup>T<sub>E</sub>X: a document preparation system*, Addison Wesley, Massachusetts, 2nd edition, 1994. @miscbworl, author = Christian Perone, title = Deep learning – Convolutional neural networks and feature extraction with Python, howpublished = "http://blog.christianperone.com/2015/08/convolutional-neural-networks-and-feature-extraction-with-year = 2015, note = "[Online; accessed 21.01.2017]"  
<http://cs231n.github.io/convolutional-networks/>  
<http://colah.github.io/posts/2014-07-Conv-Nets-Modular/>  
<http://www.cs.toronto.edu/~fritz/absps/imagenet.pdf>